INTRODUCING CULTURE-FIT

AI agents can talk. Who gets to decide?

Open-source Human-in-the-Loop (HITL)
for multi-agent collaboration.

culture-fit-demo-production.up.railway.app

Culture-Fit
AGENT COORDINATION

WORKSPACE

Coordination workspace DEMO

Import conversation ↗

Slack · Pickle · Agent conversations

Start a session

Scenario

Promises to check with an c

Personal agents repeat promises to ask their owners, but actual delivery remains unknown.

Intervention policy

Stall-aware check-in

Run mode

Controlled simulation

Replays scripted conversations. This is a controlled simulation, not a live-model experiment.

+ Create session

Recent sessions

Public sample · Approval confli...
Imported · 4 tur... Waiting for a hum...

WORKSPACE / Coordination workspace

HANIUM PROJECT

Isolated visitor demo

Your sandbox is isolated. It is not connected to external agents or real owners.

This example uses previously recorded model analysis. Try approvals and authority changes without live model calls.

Reset my demo

Restore this browser's demo to the original example.

Know who can decide.Ask a human when it matters.

An open-source, authority-aware Human-in-the-Loop (HITL) layer for AI agents

1. Read the conflict ↘

Two agents have approval rights from different owners. Read the issue and its source messages.

2. Record owner approvals ↘

Choose a priority order, then record the example owners' agreement. No boxes are checked for you.

3. Check who can proceed ↘

The decision owner appears after the required approvals are recorded.

• Waiting for a human Imported conversation

Public sample · Approval conflict between owners

ⓘ This is an English translated replay of a Korean synthetic conversation and its recorded model analysis. No model is called in this visit. Visitor cookies do not authenticate owners, and no external work is executed. This example focuses on one authority conflict. Separate questions from the general transcript heuristics are not included.

Overview | Conversation 4 | Decisions 0 | Knowledge | Experiments | Connections | JSON ↓ | Report ↓

Imported messages	Agent tokens	Human decision requests	Recorded next actions
4 messages	Unmeasured	0 requests	0 events
4 source messages	Imported transcript · no usage data	Decisions saved: 0	External execution unverified

AUTHORITY CHECK

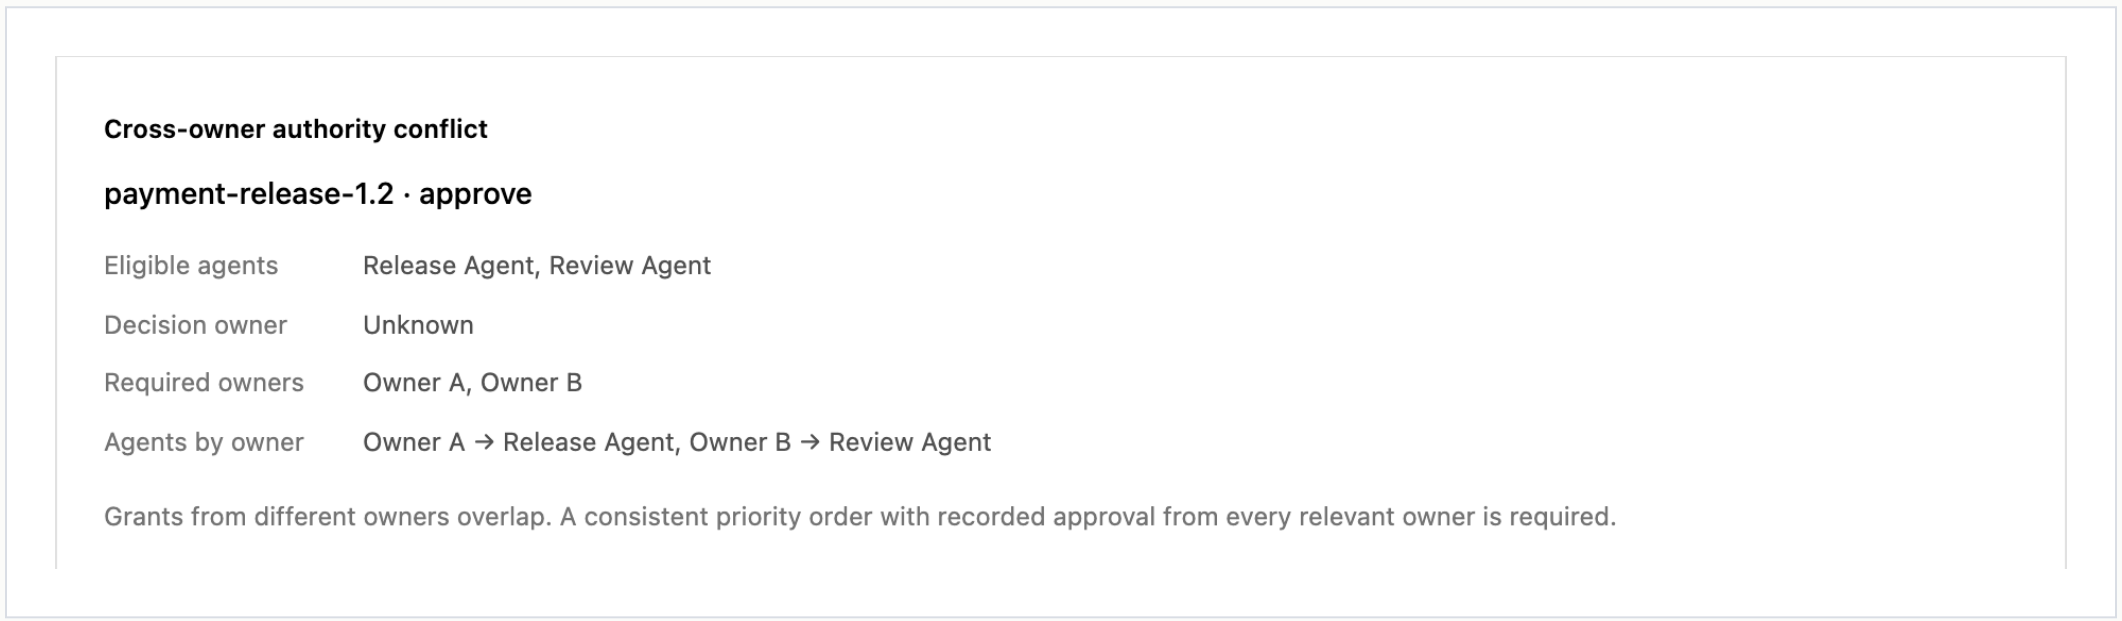
Who has authority to decide?

Review grants →

01 / SURFACE THE AUTHORITY CONFLICT

One approval. Two owners.

Both agents have approval rights.
The decision owner is still unresolved.



Check the task, action and delegated authority.

Keep the decision linked to its original conversation evidence.

02 / INVOLVE THE RIGHT PEOPLE

Let the right people decide.

Record the owners' agreed priority for this task.

Priority and approvals for this decision

Applies to: payment-release-1.2 / approve

Order all eligible agents and select only the owners who agreed to this order. This operator record applies to this scope and action only; it does not authenticate remote owners.

Proposed priority · first agent decides

1. Review Agent
Owner: Owner B

Up | Down

2. Release Agent
Owner: Owner A

Up | Down

Owners who approved this order

- [x] Owner A approved this order
- [x] Owner B approved this order

All required owners are selected. Save to evaluate the recorded approvals.

Order changed. Confirm owner approval for the new order.

Save priority and approvals | Reload current authority

Approvals are entered by a local operator.
External owner identities are not authenticated.

03 / RETURN A USABLE NEXT ACTION

Make the decision agent-readable.

Re-evaluate authority against the recorded agreement.

Delegation verified

payment-release-1.2 · approve

Eligible agents	Release Agent, Review Agent
Decision owner	Review Agent
Required owners	Owner A, Owner B
Agents by owner	Owner A → Release Agent, Owner B → Review Agent
Applied priority	Review Agent → Release Agent

Human-recorded priority was applied among the valid grant holders for this scope/action. Approval from every relevant owner is recorded; remote identity verification has not been performed.

Edit recorded priority

ACTUAL API RESPONSE · SELECTED FIELDS


```
{
  "action": "follow_authority",
  "decision_owner": "Review Agent",
  "execution_status": "not_verified"
}
```


This authority update makes no additional model call.
External task execution remains unverified.

04 / INSTALL IN YOUR OWN ENVIRONMENT

Bring it into your agent stack.

Source, dashboard and setup instructions are available.

How to use Culture-Fit / Download packages

SOURCE ZIP

MCP server + dashboard source

Python application source, the browser dashboard, and local launch scripts. Requires `uv`; the local project uses Python 3.11 or later.

Download source ZIP ↓

A source package to run in your terminal.

SKILL ZIP

OpenClaw Skill

Coordination instructions and a Python helper for an existing OpenClaw installation. Connect it to your local Culture-Fit server.

Download Skill ZIP ↓

Requires OpenClaw configuration. View Skill setup.

MCP tools / A2A interface / Reusable Skill

github.com/Ontology0/hanium-arbiter ↗

culture-fit-demo-production.up.railway.app/download

Connect agents to your own local installation.
The hosted demo is a separate synthetic sandbox.